



The Time Is Now

Introduction to SingleStore


CONFIDENTIAL

The Real-Time Distributed SQL Database





A database to
transact, analyze
and **contextualize**
data in real time.



SingleStoreDB

Multi-Model Capabilities

Relational

Standard SQL plus ability to join across the multi-model schema



Geospatial

Polygons, points, lines and built-in spatial functions



Unstructured, full-text search

Built-in, native inverted indexing



Key-value

Scan >30B tuple space and return trillions of key-value results in <1 second



JSON/document

JSON as native type for schema-on-write and schema-on-read



Time-series

Native date, time, and timestamp types with built-in functions, e.g. time_bucket



Vector functions

dot_product and euclidean distance for nearest neighbour and semantic search



SingleStoreDB

Key product capabilities

Universal Storage

Patented single table type for transactions and analytics. Compatibility with ANSI SQL, MySQL + MariaDB ecosystem; NoSQL



Pure speed

Fast transactions, analytics, vectorization and query compilation



Fast ingestion

SingleStore Pipelines — load data w/updates



Unlimited storage

Separation of storage + compute in a unified database



MongoDB® API

SingleStore Kai™ powers up to 1,000x faster analytics on JSON



Multi-model

Relational foundation with support for various data types, including vectors



Run anywhere


Hybrid, multi-cloud, SaaS, on-premises, Kubernetes operator




Intelligence

Build generative AI and other apps with **Notebooks** on top of SingleStoreDB






SingleStoreDB: Under the Hood




Pipelines: Real-Time Ingest

Product overview

- Millions of records / sec with immediate availability
- Built-in component of database
- Transactional consistency
- Exactly-once semantics
- Native integrations with Kafka, cloud storage, HDFS
- Ingest to stored procedures to shape/modify data
- Pipelines support JSON, Avro, Parquet, and CSV data formats.




SingleStoreDB as Vector Database



Using SingleStoreDB for your AI

Experience faster vectors with an exact next neighbor match and better compression. It's an enterprise grade solution that powers fast ingestion with ms response times (no ETL), and enables joins with your app data.



Adopting SingleStore

1

Build new modern applications

Net new apps or operational analytics workloads



2

Augment Tx or analytical DBs

Augment legacy and speciality OLTP/ OLAP engines, or Hadoop



3

Consolidate and modernize

Modernize or replace legacy datastores, speciality databases, or Hadoop

